

CSCI 3907/6907

Fall 2019

Lecture 2 – Part I

Regular Expressions & Automata

Notes

- Syllabus: *Check Often!*

<https://sites.google.com/view/csci-39076907-fall2019/syllabus>

- four main assignments 40%, Midterm 30%, Final Project 30%

- Register for Piazza

<https://piazza.com/gwu/fall2019/csci39076907nlp/home>

- Start forming a group for final project (max 4 students) and choosing your topic
- The first assignment will be released this week.
 - *All assignment related questions must be directed to the TA.*
 - *You have 7 free grace days for the whole semester.*

Review

- *What is NLP?*
 - *Building programs that can recognize, analyze , and generate text and speech*
 - *Processing unstructured data*
- NLP Applications: Machine Translation, Sentiment Analysis, Part of Speech Tagging ... etc
- Ambiguity results from the existence of multiple possibilities for each level
- NLP Approaches: Rule based and statistical based approaches

Outline: Basic Text Processing


- Part 1: Formal languages: regular Expressions and automata
- Part 2: Words and Transducers

Formal Languages

- A **formal language** is a set of **strings**, each string composed of symbols from a finite symbol-set called an **alphabet** Σ .
- A model which can both **generate** and **recognize** all and only the strings of a formal language acts as a definition of the formal language.

Formal Languages – Cont.

- A formal language may bear no resemblance at all to a real language (natural language)
- but can model a subset of expressions:
 - Parts of the phonology, morphology, or syntax.
 - Date and time expressions.
- Example: Sheeptalk
 - Any string from the following (infinite) set:
 - Baa!
 - Baaa!
 - Baaaa!
 - Baaaaa!
 -



Regular Languages

- A regular language \rightarrow a formal language that can be fully specified using one of the following:
 - Regular expression
 - Finite state automaton
 - Regular grammar

Regular Expressions

- A formal language for specifying text search strings
- Used in
 - Unix tools like grep and sed
 - Programming languages: Perl, python, Java, etc.

Regular Expressions

- Requires:
 - a pattern (e.g. all words that start with *ca*)
 - a corpus of texts to search through
- Returns all text that matches the pattern or only the first match

Regular Expressions: Disjunctions

- Letters inside square brackets []

Pattern	Matches
[wW]oodchuck	Woodchuck, woodchuck
[1234567890]	Any digit

- Ranges [A-Z]

Pattern	Matches	Examples
[A-Z]	An upper case letter	<u>D</u> renched Blossoms
[a-z]	A lower case letter	<u>m</u> y beans were impatient
[0-9]	A single digit	Chapter <u>1</u> : Down the Rabbit Hole

Negation in Disjunction

- Negations [^Ss]
 - Carat means negation only when first in []

Pattern	Matches	Examples
[^A-Z]	Not an upper case letter	O <u>y</u> fn pripetchik
[^Ss]	Neither 'S' nor 's'	I have no exquisite reason"
[^e^]	Neither e nor ^	<u>L</u> ook here
a^b	The pattern a carat b	Look up <u>a^b</u> now


More Disjunction

- Woodchucks is another name for groundhog!
- The pipe | for disjunction

Pattern	Matches
groundhog woodchuck	groundhog, woodchuck
yours mine	yours , mine
a b c	a, b, c
[gG]roundhog [Ww]oodchuck	groundhog, Groundhog, woodchuck, Woodchuck

? * + .

Pattern	Matches	
colou?r	Optional previous char	<u>color</u> <u>colour</u>
oo*h!	0 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
o+h!	1 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
baa+		<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
beg.n		<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>



Stephen C Kleene

Anchors ^ \$

Pattern	Matches
<code>^[A-Z]</code>	<u>P</u> alo Alto
<code>^[^A-Za-z]</code>	<u>1</u> <u>“Hello”</u>
<code>\.\$</code>	The end <u>.</u>
<code>.\$</code>	The end <u>?</u> The end <u>!</u>

Other Operators & Special Characters

Pattern	matches
\d	[0-9], any digit
\D	Any non digit
\w	Any alphanumeric or space
\W	[^\w] any non-alphanumeric
\s	whitespace
\S	non whitespace
*, \?, \.	*, ?, . (Characters that need to be escaped)
\n	end of line
\t	tab
{n}	n occurrences of the previous char or expression
{n,m}	n to m occurrences of the previous char or expression
{n, }	At least n occurrences of the previous char or expression

Substitution And Memory

- Find a pattern in text, then replace it with another pattern.
`s/pattern/replacement/`
- Capture groups: store a matched pattern in memory, then reuse it in substitution or later matches:
 - The resulting match is stored in a numbered register (e.g `\1` and `\2`)
 - Using memory in search:
`/the (.*)er they were, the \1er they will be/`
`/the (.*)er they (.*) , the \1er we \2/`
 - Using memory in substitution: `s/([0-9]+)/<\1>/`

ELIZA

- Use regular expressions and capture groups to generate responses:

User: You are like my father in some ways.

ELIZA: WHAT RESEMBLANCE DO YOU SEE

User: You are not very aggressive but I think you don't want me to notice that.

ELIZA: WHAT MAKES YOU THINK I AM NOT VERY AGGRESSIVE

User: You don't argue with me.

ELIZA: WHY DO YOU THINK I DON'T ARGUE WITH YOU

User: You are afraid of me.

ELIZA: DOES IT PLEASE YOU TO BELIEVE I AM AFRAID OF YOU

Weizenbaum (1966)

```
s/. * I'M (depressed|sad) . */I AM SORRY TO HEAR YOU ARE \1/
s/. * I AM (depressed|sad) . */WHY DO YOU THINK YOU ARE \1/
s/. * all . */IN WHAT WAY/
s/. * always . */CAN YOU THINK OF A SPECIFIC EXAMPLE/
```

Simple ELIZA Implementation

#Upper case all characters

tr/a-z/A-Z/;

#Remove all non-words except for spaces

s/[^\w\s]//;

#Switch first person to second person

s/\bMY\b/YOUR/g;

s/\b(I|ME)\b/YOU/g;

s/\bAM\b/ARE/g;

#Responses

s/.*\b(HI|HELLO|HOWDY)\b.*/\$1 THERE. HOW ARE YOU DOING?/;

s/.*\bYOU ARE (DEPRESSED|SAD|UNHAPPY)\b.*/I AM SORRY TO HEAR YOU ARE \$1./;

s/.*\bALL\b.*/IN WHAT WAY?/;

s/.*\bALWAYS\b.*/CAN YOU THINK OF A SPECIFIC EXAMPLE?/;

s/.*\bNOT\b.*/WHY NOT?/;

s/.*\bYES\b.*/I SEE/;

#Add prompt

s/^/>> /;

Example

- Find me all instances of the word “the” in a text.

the

e.g. misses capitalized examples

[tT]he

e.g. incorrectly returns *other* or *theology*

[^a-zA-Z][Tt]he[^a-zA-Z]

Errors

- Two kinds of errors:
 - False positive (Type I)
 - Matching strings that we should not have matched (**there**, **then**, **other**)
 - False negative (Type II)
 - Not matching things that we should have matched (The)

Errors – Cont.


- We always deal with these types of errors in NLP
- Reduce the error rate for an application involve:
 - Increasing accuracy or precision (*minimizing false positives*)
 - Increasing coverage or recall (*minimizing false negative*)

Summary

- Regular expressions play a surprisingly large role
 - Sophisticated sequences of regular expressions are often the first model for any text processing text
 - Early version of chatbots (ELIZA), text normalization, tokenization, stemming, sentence segmentation, string similarities
- For many hard tasks, we use machine learning classifiers
 - But regular expressions are used as features in the classifiers
 - Can be very useful in capturing generalizations

Finite State Automata (FSA)



- Abstract model of computation
- Consists of N **states** and a **transition** matrix.
- A **start** state, and a set of final “**accept**” states
- Can be deterministic or non deterministic



	Input		
State	b	a	!
0	1	∅	∅
1	∅	2	∅
2	∅	3	∅
3	∅	3	4
4:	∅	∅	∅


Formally

- A regular language is a 5-tuple consisting of
 - Q : set of states $\{q_0, q_1, q_2, q_3, q_4\}$
 - Σ : an alphabet of symbols $\{a, b, !\}$
 - q_0 : a start state in Q
 - F : a set of final states in Q $\{q_4\}$
 - $\delta(q, i)$: a transition function mapping $Q \times \Sigma$ to Q





Finite State Automata (FSA)

- Adding a **fail** state



The Tape Metaphor

- Imagine the input being written on a long tape broken into cells:



Deterministic FSAs

- Deterministic: at each point in processing there is always one unique thing to do (no choices).
- **D-RECOGNIZE** is a simple table-driven interpreter
 - The algorithm is universal for all unambiguous languages.
 - To change the machine, you change the table.
- FSAs can be useful tools for recognizing – and generating – subsets of natural language

Deterministic Recognition Algorithm

function D-RECOGNIZE(*tape*, *machine*) **returns** accept or reject

index \leftarrow Beginning of tape

current-state \leftarrow Initial state of machine

loop

if End of input has been reached **then**

if *current-state* is an accept state **then**

return accept

else

return reject

elseif *transition-table*[*current-state*, *tape*[*index*]] is empty **then**

return reject

else


current-state \leftarrow *transition-table*[*current-state*, *tape*[*index*]]

index \leftarrow *index* + 1


end

Non-Deterministic FSAs

- Deterministic sheep language FSA



- Non-deterministic sheep language FSA



Dealing with non-determinism

- At any choice point, we may follow the wrong arc
- Potential solutions:
 - Save **backup** states at each choice point
 - **Look-ahead** in the input before making choice
 - Pursue alternatives in **parallel**

Equivalence

- For every NFSA there exists an equivalent DFSA
 - (i.e. that accepts exactly the same set of strings).
- The resulting DFSA, may have many more states than the original NFSA
 - (up to 2^N states for a NFSA with N states).